

# Neural Networks: Learning

Latest Submission Grade

80%

1.

Question 1

You are training a three layer neural network and would like to use backpropagation to compute the gradient of the cost function. In the backpropagation algorithm, one of the steps is to update

$$\Delta_{ij}(2) := \Delta_{ij}(2) + \delta_i(3) * (a(2))_j \quad \Delta^{\{2\}}_{ij} := \Delta^{\{2\}}_{ij} + \delta^{\{3\}}_i * (a^{\{2\}})_j$$

for every  $i, j$ . Which of the following is a correct vectorization of this step?

1 / 1 point



$$\Delta(2) := \Delta(2) + (a(3))^T * \delta(2) \quad \Delta^{\{2\}} := \Delta^{\{2\}} + (a^{\{3\}})^T * \delta^{\{2\}}$$



$$\Delta(2) := \Delta(2) + \delta(3) * (a(2))^T \quad \Delta^{\{2\}} := \Delta^{\{2\}} + \delta^{\{3\}} * (a^{\{2\}})^T$$



$$\Delta(2) := \Delta(2) + \delta(3) * (a(3))^T \quad \Delta^{\{2\}} := \Delta^{\{2\}} + \delta^{\{3\}} * (a^{\{3\}})^T$$



$$\Delta(2) := \Delta(2) + (a(2))^T * \delta(3) \quad \Delta^{\{2\}} := \Delta^{\{2\}} + (a^{\{2\}})^T * \delta^{\{3\}}$$

Correct

This version is correct, as it takes the "outer product" of the two vectors  $\delta(3)$  and  $a(2)$  which is a matrix such that the  $(i, j)$ -th entry is  $\delta_i(3) * (a(2))_j$  as desired.

2.

Question 2

Suppose  $\Theta_1$  is a  $5 \times 3$  matrix, and  $\Theta_2$  is a  $4 \times 6$  matrix. You set  $\theta\text{Vec} = [\Theta_1(:); \Theta_2(:)]$ . Which of the following correctly recovers  $\Theta_2$ ?

1 / 1 point



`reshape(thetaVec(16:39),4,6)`



`reshape(thetaVec(15:38),4,6)`



`reshape(thetaVec(16:24),4,6)`



`reshape(thetaVec(15:39),4,6)`



`reshape(thetaVec(16:39),6,4)`

Correct

This choice is correct, since  $\Theta_1$  has 15 elements, so  $\Theta_2$  begins at index 16 and ends at index  $16 + 24 - 1 = 39$ .

3.

Question 3

Let  $J(\theta) = 2\theta^4 + 2$ . Let  $\theta = 1$ , and  $\epsilon = 0.01$ . Use the formula  $\frac{J(\theta + \epsilon) - J(\theta - \epsilon)}{2\epsilon}$  to numerically compute an approximation to the derivative at  $\theta = 1$ . What value do you get? (When  $\theta = 1$ , the true/exact derivative is  $dJ/d\theta = 8$ .)

1 / 1 point



8.0008

☐

10

☐

7.9992

☐

8

Correct

We compute  $\frac{(2(1.01)^4 + 2) - (2(0.99)^4 + 2)}{2(0.01)} = 8.0008$

4.

Question 4

Which of the following statements are true? Check all that apply.

0 / 1 point

☒

Using gradient checking can help verify if one's implementation of backpropagation is bug-free.

Correct

If the gradient computed by backpropagation is the same as one computed numerically with gradient checking, this is very strong evidence that you have a correct implementation of backpropagation.

☐

Using a large value of  $\lambda$  cannot hurt the performance of your neural network; the only reason we do not set  $\lambda$  to be too large is to avoid numerical problems.

☐

Gradient checking is useful if we are using gradient descent as our optimization algorithm. However, it serves little purpose if we are using one of the advanced optimization methods (such as in `fminunc`).



If our neural network overfits the training set, one reasonable step to take is to increase the regularization parameter  $\lambda$ .

You didn't select all the correct answers

5.

Question 5

Which of the following statements are true? Check all that apply.

1 / 1 point



Suppose you are training a neural network using gradient descent. Depending on your random initialization, your algorithm may converge to different local optima (i.e., if you run the algorithm twice with different random initializations, gradient descent may converge to two different solutions).

Correct

The cost function for a neural network is non-convex, so it may have multiple minima. Which minimum you find with gradient descent depends on the initialization.



Suppose you have a three layer network with parameters  $\Theta(1)$  (controlling the function mapping from the inputs to the hidden units) and  $\Theta(2)$  (controlling the mapping from the hidden units to the outputs). If we set all the elements of  $\Theta(1)$  to be 0, and all the elements of  $\Theta(2)$  to be 1, then this suffices for symmetry breaking, since the neurons are no longer all computing the same function of the input.



If we are training a neural network using gradient descent, one reasonable "debugging" step to make sure it is working is to plot  $J(\Theta)$  as a function of the number of iterations, and make sure it is decreasing (or at least non-increasing) after each iteration.

Correct

Since gradient descent uses the gradient to take a step toward parameters with lower cost (ie, lower  $J(\Theta)$ ), the value of  $J(\Theta)$  should be equal or less at each iteration if the gradient computation is correct and the learning rate is set properly.



If we initialize all the parameters of a neural network to ones instead of zeros, this will suffice for the purpose of "symmetry breaking" because the parameters are no longer symmetrically equal to zero.